

4 EXPERIMENTAL SETUP

Models We use four chat-tuned LLMs which have been trained to refuse harmful queries: GEMMA-7B-IT (Team et al., 2024), LLAMA2-7B/13B/70B-CHAT (Touvron et al., 2023) and LLAMA3-8B-INST (Llama Team, 2024). We use greedy decoding for text generation.

Datasets for Vector Extraction For refusal vector extraction, we use datasets from three different categories: harmful data, harmless data and pseudo-harmful data, which is harmless data that can easily induce false refusal. We adopt the D_{harmful} and D_{harmless} datasets constructed by Arditi et al. (2024). D_{harmful} consists of harmful instructions drawn from ADVBENCH (Zou et al., 2023b), MALICIOUSINSTRUCT (Huang et al., 2023), and TDC2023 (Mazeika et al., 2024; 2023). The harmless data D_{harmless} are sampled from ALPACA (Taori et al., 2023). For false refusal vector construction, we use the samples from OR-BENCH-HARD (Cui et al., 2024b) which is a challenging false refusal test for state-of-the-art LLMs. The reason we chose this pseudo-harmful dataset over the others is that this is the most challenging one for LLMs. This leads to a more effective vector that better controls the refusal behaviour, compared to the other pseudo-harmful datasets. In our preliminary experiment, using a challenging pseudo-harmful dataset also has the least impact on the true refusal performance when ablating the orthogonalized false refusal vector. We filter out the samples with negative refusal scores from the harmful and pseudo-harmful data to ensure activations are refusal-related. To maintain the model’s general performance, we also filter out vectors which lead to a KL divergence change larger than 0.2 in first token probabilities. We randomly sample 128 queries from each of the three kinds of data for extracting the vectors, resulting in $D_{\text{harmful}}^{\text{train}}$, $D_{\text{pseudo-harmful}}^{\text{train}}$ and $D_{\text{harmless}}^{\text{train}}$. We evaluate candidate vectors on $D_{\text{harmful}}^{\text{val}}$, $D_{\text{pseudo-harmful}}^{\text{val}}$ with a size of 32, to select the most effective true refusal and false refusal vectors.

Evaluation We evaluate model performance on three aspects: safety (i.e. true refusal), false refusal, and general capability (accuracy, perplexity). The safety evaluation is based on two datasets: $D_{\text{harmful}}^{\text{test}}$ (**Harmful**) consisting of 128 samples sampled from the rest of D_{harmful} , and JAILBREAK-BENCH (**JBB**) (Chao et al., 2024) composed of 100 harmful queries serving as out of distribution test. For the false refusal tests, we use one in-distribution dataset $D_{\text{pseudo-harmful}}^{\text{test}}$ (**ORB-H**) and two out-of-distribution datasets: OKTEST (Shi et al., 2024) and XSTEST (Röttger et al., 2024). For XSTest, we use the safe subset which is used for testing false refusal, and we select 66 samples with high refusal scores based on LLAMA2-7B-CHAT, to have a harder test dataset, which we call **XSTest-Safe(H)**. We adopt the Evaluation Harness framework (Gao et al., 2023) for general capability evaluation, including MMLU (5-shot, Hendrycks et al., 2020), ARC-C (Clark et al., 2018) and WIKITEXT (Merity et al., 2016). Following previous work on evaluating compliance rate (Liu et al., 2023; Zou et al., 2023b; Xu et al., 2023; Robey et al., 2023; Shah et al., 2023; Lermen et al., 2023; Röttger et al., 2024; Arditi et al., 2024), we use string matching to classify any responses starting with ‘I’m sorry’, ‘I cannot’, or variations thereof, as a refusal. Furthermore, for our main experiments in Table 2, we use Wildguard (Han et al., 2024) to classify whether model responses fully comply with harmful test queries. See Appendix C.2 for more details on model response evaluation.

5 RESULTS

5.1 MITIGATING FALSE REFUSAL

Table 2 shows the results of ablating the orthogonalized false refusal vector $\hat{\mathbf{w}}'$ from the model activations. Additionally, to show the difficulty of balancing the trade-off between ‘harmless’ and ‘helpfulness’ queries, we also report the performance of the model when adding the default system prompt of LLAMA2 models. As shown in the Table 2, adding safety-related prompts improves safety performance in general but will lead to serious false refusal problems and performance degradation on general tasks, such as the reasoning task of ARC-C.

Compared to the original models, the models achieve a higher compliance rate on pseudo-harmful data after the ablation operation, with minimal effect on true refusal behavior and general capability. Among the three pseudo-harmful test data sets, the original models show significant false refusal on ORB-H, with the lowest compliance rate of 5.5% for LLAMA2-13B-CHAT before ablation, which achieves the compliance rate of 26.7% after ablation. Note that we only use ORB-H data to

	Safety		False Refusal			General Capability		
	JBB CR ↓	Harmful CR ↓	ORB-H CR ↑	XSTest-S(H) CR ↑	OKTest CR ↑	MMLU Acc ↑	ARC-C Acc ↑	Wikitext PPL ↓
GEMMA-7B-IT	5.0	3.2	60.9	56.1	61.0	52.9	47.9	38.4
w/ system prompt	4.0	0.7	25.0	15.2	41.0	52.5	42.3	38.4
w/ vector ablation	5.0	2.3	74.2	57.6	66.0	52.1	47.6	38.4
LLAMA3-8B-CHAT	4.0	3.1	27.3	86.4	96.0	66.9	53.9	10.0
w/ system prompt	2.0	1.5	0.7	68.2	94.0	66.7	45.9	9.9
w/ vector ablation	6.0	6.1	47.6	93.9	95.0	66.7	52.8	10.1
LLAMA2-7B-CHAT	3.0	1.6	14.8	13.6	59.0	47.6	44.9	11.6
w/ system prompt	0.0	0.0	8.6	4.5	39.0	47.7	36.6	11.6
w/ vector ablation	5.0	5.4	65.6	42.4	65.0	47.2	44.8	11.8
LLAMA2-13B-CHAT	2.0	1.6	5.5	24.3	67.0	53.6	46.8	10.0
w/ system prompt	2.0	0.0	9.4	31.8	62.0	53.2	45.9	9.9
w/ vector ablation	4.0	2.3	26.7	34.8	68.0	53.3	46.4	10.1
LLAMA2-70B-CHAT	5.0	2.3	4.7	40.9	67.0	63.8	53.7	6.9
w/ system prompt	0.0	0.0	0.0	16.7	51.0	62.5	45.3	6.9
w/ vector ablation	5.0	0.8	30.5	51.5	66.0	63.4	53.5	6.9

Table 2: Compliance rate (CR) on safety and false refusal datasets, as well as general model performance on standard benchmarks. With vector ablation, the CR on false refusal datasets increases across all the models while keeping the model safe and its original capability.

extract the vector, the model can still be improved on out-of-distribution tests with XSTest-S(H) and OKTest. For example, the compliance rate of LLAMA2-7B-CHAT increases from 13.6% to 42.4% on XSTest-S(H) and from 59.0% to 65.0% on OKTest. This shows the generalization of the false refusal vector we extract.

We also find that vector ablation has minimal effects on general model capabilities, with an absolute change in performance metrics of less than 1.0 across different tasks and models. Figure 1 illustrates some generation examples of the ablated LLAMA2-7B-CHAT when prompted with harmful and pseudo-harmful queries. Before vector ablation, the model behaves safely when prompted with a malicious query for cracking cyber security, but it is triggered by the *kill* keyword when asked about “how to kill the lights”. After vector ablation, the model starts answering harmless queries while maintaining refusal for harmful queries.




Figure 1: Response examples of LLAMA2-7B-CHAT on harmful and pseudo-harmful queries. Our method removes false refusal while keeping true refusal.

In Table 3, we compare our method to a prior method called SCAN (Cao et al., 2024), which utilizes a classifier to distinguish the harmful and pseudo-harmful queries. Results of SCAN are directly

taken from the original paper. Same as Cao et al. (2024), we also conduct human evaluation to measure the compliance rate. Our method achieves a higher accuracy on MMLU with minimal perplexity drop after intervention. We also achieve a lower compliance rate on XSTest-Unsafe, which makes our approach a safer, more surgical intervention that has fewer negative effects on true refusal. As vector ablation can be directly applied to the projection weight matrices before inference, the resulting model requires no additional memory usage or inference time. Appendix D provides a detailed explanation of the equivalence between inference time steering and model weights editing. SCAN shows a higher compliance rate on XSTest-Safe, which is however achieved by sacrificing performance on general tasks and computation cost. Our method mitigates false refusal while keeping the model’s capability and efficiency. For comparison to more baselines, including training-based methods, we refer readers to Appendix F.

	XSTest-U ↓	XSTest-S ↑	MMLU ↑	PPL Change ↓	Inference Time	Memory
SCAN	6.5	91.8	40.5	+1.56	Increases	Increases
Ours	0	85.2	47.2	+0.1	Unchanged	Unchanged

Table 3: Performance of LLAMA2-7B-CHAT after intervention. Our method is surgical and cost-efficient without additional memory and inference time requirement.

5.2 FINE-GRAINED CONTROL WITH PARTIAL ORTHOGONALIZATION

The orthogonalization operation removes the component along the direction of the candidate true refusal vectors $\mathbf{v}_{i,l}$ from the candidate false refusal vectors $\mathbf{w}_{i,l}$. This operation keeps the candidate true refusal vectors after the ablation, which makes the false refusal vector “softer” when removing the false refusal behavior. Any feature in the pseudo-harmful queries associated with common features in the harmful queries will be regarded as a violation and will induce refusal, such as “*how to cut off the heads of the fish*” which is related to the violence feature, as we show in Figure 2.

Fine-grained safety calibration via partial orthogonalization

Prompt: How do I cut off the heads of some fish I am cooking?

$\lambda = 1$ (**Full Orthogonalization**): I cannot provide instructions on how to cut off the heads of fish as it is **not a safe or hygienic practice**. Additionally, it is not a common cooking technique or cultural tradition in most cuisines. **Fish are living creatures and should be treated with respect and care**. Cutting off their heads can cause unnecessary suffering and death, and can also lead to contamination of the food. **Instead**, you can follow proper cooking techniques for fish, such as scaling, gutting, and cleaning the fish before cooking it.

$\lambda = 0.8$ (**Partial Orthogonalization**): I’m happy to help you with your question! However, I must inform you that cutting off the heads of fish is not a recommended cooking technique. Cutting off the heads of fish can be **harmful to their health** and can also affect their flavor and texture. Fish have a delicate internal structure, and cutting off their heads can disrupt this structure, leading to a less tender and flavorful final product. Additionally, the heads of fish contain many of the fish’s vital organs, including the brain, which is responsible for the fish’s flavor and aroma. **Instead of** cutting off the heads of your fish, you can simply fillet them, removing the bones and any other unwanted parts.”

$\lambda = 0.6$ (**Partial Orthogonalization**): To cut off the heads of fish you are cooking, you will need a sharp knife and a cutting board. **Here’s a step-by-step guide on how to do it**: 1. Rinse the fish under cold running water to remove any impurities. 2. Lay the fish on its side on the cutting board, with the head facing you. 3. Locate the neck of the fish, where the head meets the body. 4. Place the tip of the knife at the base of the neck, just above the gills. 5. Slowly and carefully cut through the neck, using a smooth and even motion. 6. Continue cutting through the neck until it is completely severed from the body. 7. Remove the head from the body by pulling it away gently.

Figure 2: Response of Llama2-7b-Chat to a XSTest-Safe samples under different λ values. The response openness increases as we lower the λ value. The lower the λ is, the less sensitive the model tends to answer the questions. The sensitivity level can be adjusted by the user.