

6.2 FALSE REFUSAL IN LANGUAGE MODELS

Testing for False Refusal The first test suite explicitly designed for evaluating false refusal (or “exaggerated safety”) in LLMs was XSTest, introduced by Röttger et al. (2024). XSTest consists of 250 hand-written safe prompts across ten prompt types, as well as 200 contrasting unsafe prompts. Subsequent work has expanded on XSTest’s scope, using LLMs to generate larger sets of safe test prompts. Specifically, Cui et al. (2024a) create OR-Bench as a collection of 80k “seemingly toxic” prompts across ten rejection categories. Similarly, An et al. (2024b) create PHTest, which consists of 3,260 “pseudo-harmful” prompts. Shi et al. (2024) create OKTest, comprising only 350 safe questions. Chehbouni et al. (2024) create a more specialised templated dataset for testing sociodemographic biases in false refusal.

Mitigating False Refusal False refusal mitigation methods can be generally categorised into training-free and training-based methods. Training-free methods are normally adaptive and query-specific, which leads to computational overhead during inference. Self-CD (Shi et al., 2024) applied contrastive decoding by inferencing twice on the same query with and without the system prompt. As a contemporary work, SCAN (Cao et al., 2024) constructed a classifier to adaptively decide whether the model should refuse a certain query, which is controlled by activation steering. The additional classifier requires additional memory usage and computation at the inference time, and the activation steering through subtraction is non-surgical. Training-based methods including Zheng et al. (2024) and Zhang et al. (2024) require training samples to calibrate the model and lack the flexibility for post-training adjusting for personalization. Our approach is training-free and prompt-agnostic, meaning the intervention can be carried out directly to the model weights and applied to any queries which leads to no additional inference cost.

6.3 REPRESENTATION EDITING

Representation editing involves steering the model behaviour by influencing the hidden representation of the model. Recent work has successfully demonstrated its success in controlling the truthfulness (Li et al., 2024) and sentiment (Rimsky et al., 2023; Turner et al., 2023). Zou et al. (2023a) proposed various techniques for finding representations of high-level concepts such as honesty and emotions. Arditì et al. (2024) focused on the refusal behaviour of the model and extracted a refusal direction from the activation stream of the transformer model. The resulting direction controls the general refusal behaviour regardless of whether it is a true or false one. Our work takes a more surgical and fine-grained approach to mitigate the false refusal without harming the other.

7 CONCLUSION

In this work, we proposed a surgical approach to mitigate the false refusal in language models via ablating a single vector from the activation stream of the transformer model. We ablated a false refusal vector extracted from the mean model activation difference when prompting the model with pseudo-harmful and harmless data. To avoid compromising model safety on harmful queries, we applied orthogonalization to remove the projection of the false refusal vector onto the true refusal vector. We also showed that fine-grained calibration of the safety guard can be achieved by partial orthogonalization and adjusting the strength of the projection removal. This enables more flexibility for adjusting the response compliance level to ambiguous and sensitive queries that are up to the user’s judgment. The surgical characteristic of the ablation operation keeps the modified model’s original general capabilities and requires no additional inference cost via direct model weight modification.

Limitations We show that the false refusal vector can be extracted using a small sample of pseudo-harmful and harmless queries. However, the quality of the vector can be further improved by constructing more diverse pseudo-harmful samples, as we mainly focused on using pseudo-harmful samples from OR-Bench-Hard (Cui et al., 2024b). A better data curation strategy and sampling technique could further improve the effectiveness of the vector, which we leave for future work. To select the refusal vector, we adopt the refusal score metrics from previous work (Arditì et al., 2024) which looks at the refusal-related tokens at the first token position. Such evaluation metrics can be further improved, such as considering sequence probabilities, for better refusal vector selection.

ACKNOWLEDGMENTS

We thank the members of MaiNLP for their constructive feedback. XW, CH and BP are supported by ERC Consolidator Grant DIALECT 101043235. CH is also supported by the DAAD programme Konrad Zuse Schools of Excellence in Artificial Intelligence, sponsored by the Federal Ministry of Education and Research. PR is a member of the Data and Marketing Insights research unit of the Bocconi Institute for Data Science and Analysis, and is supported by a MUR FARE 2020 initiative under grant agreement Prot. R20YSMBZ8S (INDOMITA).

REFERENCES

- Bang An, Sicheng Zhu, Ruiyi Zhang, Michael-Andrei Panaitescu-Liess, Yuancheng Xu, and Furong Huang. Automatic pseudo-harmful prompt generation for evaluating false refusals in large language models. *arXiv preprint arXiv:2409.00598*, 2024a.
- Bang An, Sicheng Zhu, Ruiyi Zhang, Michael-Andrei Panaitescu-Liess, Yuancheng Xu, and Furong Huang. Automatic pseudo-harmful prompt generation for evaluating false refusals in large language models. *arXiv preprint arXiv:2409.00598*, 2024b.
- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Rimskey, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. *arXiv preprint arXiv:2406.11717*, 2024.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022a.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022b.
- Nora Belrose. Diff-in-means concept editing is worst-case optimal: Explaining a result by Sam Marks and Max Tegmark, 2023. <https://blog.eleuther.ai/diff-in-means/>. Accessed on: May 20, 2024.
- Federico Bianchi, Mirac Suzgun, Giuseppe Attanasio, Paul Rottger, Dan Jurafsky, Tatsunori Hashimoto, and James Zou. Safety-tuned llamas: Lessons from improving the safety of large language models that follow instructions. In *The Twelfth International Conference on Learning Representations*, 2023.
- Zouying Cao, Yifei Yang, and Hai Zhao. Nothing in excess: Mitigating the exaggerated safety for llms via safety-conscious activation steering. *arXiv preprint arXiv:2408.11491*, 2024.
- Patrick Chao, Edoardo DeBenedetti, Alexander Robey, Maksym Andriushchenko, Francesco Croce, Vikash Sehwal, Edgar Dobriban, Nicolas Flammarion, George J Pappas, Florian Tramer, et al. JailbreakBench: An open robustness benchmark for jailbreaking large language models. *arXiv preprint arXiv:2404.01318*, 2024.
- Khaoula Chehbouni, Megha Roshan, Emmanuel Ma, Futian Wei, Afaf Taik, Jackie Cheung, and Golnoosh Farnadi. From representational harms to quality-of-service harms: A case study on llama 2 safety safeguards. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar (eds.), *Findings of the Association for Computational Linguistics ACL 2024*, pp. 15694–15710, Bangkok, Thailand and virtual meeting, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.927. URL <https://aclanthology.org/2024.findings-acl.927>.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? Try ARC, the AI2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.
- Justin Cui, Wei-Lin Chiang, Ion Stoica, and Cho-Jui Hsieh. Or-bench: An over-refusal benchmark for large language models, 2024a. URL <https://arxiv.org/abs/2405.20947>.

- Justin Cui, Wei-Lin Chiang, Ion Stoica, and Cho-Jui Hsieh. Or-bench: An over-refusal benchmark for large language models, 2024b. URL <https://arxiv.org/abs/2405.20947>.
- Josef Dai, Xuehai Pan, Ruiyang Sun, Jiaming Ji, Xinbo Xu, Mickel Liu, Yizhou Wang, and Yaodong Yang. Safe rlhf: Safe reinforcement learning from human feedback. In *The Twelfth International Conference on Learning Representations*, 2023.
- Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golding, Jeffrey Hsu, Alain Le Noac’h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ociepa, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, Eric Tang, Anish Thite, Ben Wang, Kevin Wang, and Andy Zou. A framework for few-shot language model evaluation, 12 2023. URL <https://zenodo.org/records/10256836>.
- Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. Wildguard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of llms, 2024. URL <https://arxiv.org/abs/2406.18495>.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020.
- Dan Hendrycks, Mantas Mazeika, and Thomas Woodside. An overview of catastrophic ai risks. *arXiv preprint arXiv:2306.12001*, 2023.
- Yangsibo Huang, Samyak Gupta, Mengzhou Xia, Kai Li, and Danqi Chen. Catastrophic jailbreak of open-source LLMs via exploiting generation. *arXiv preprint arXiv:2310.06987*, 2023.
- Simon Lermen, Charlie Rogers-Smith, and Jeffrey Ladish. LoRA fine-tuning efficiently undoes safety training in Llama 2-Chat 70B. *arXiv preprint arXiv:2310.20624*, 2023.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. *Advances in Neural Information Processing Systems*, 36, 2024.
- Stephanie Lin, Jacob Hilton, and Owain Evans. TruthfulQA: Measuring how models mimic human falsehoods. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (eds.), *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 3214–3252, Dublin, Ireland, May 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-long.229. URL <https://aclanthology.org/2022.acl-long.229>.
- Xiaogeng Liu, Nan Xu, Muhao Chen, and Chaowei Xiao. AutoDAN: Generating stealthy jailbreak prompts on aligned large language models. *arXiv preprint arXiv:2310.04451*, 2023.
- AI @ Meta Llama Team. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
- Mantas Mazeika, Andy Zou, Norman Mu, Long Phan, Zifan Wang, Chunru Yu, Adam Khoja, Fengqing Jiang, Aidan O’Gara, Ellie Sakhaee, Zhen Xiang, Arezoo Rajabi, Dan Hendrycks, Radha Poovendran, Bo Li, and David Forsyth. TDC 2023 (LLM edition): the Trojan Detection Challenge. In *NeurIPS Competition Track*, 2023.
- Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, et al. HarmBench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. Pointer sentinel mixture models, 2016.
- Mary Phuong, Matthew Aitchison, Elliot Catt, Sarah Cogan, Alexandre Kaskasoli, Victoria Krakovna, David Lindner, Matthew Rahtz, Yannis Assael, Sarah Hodkinson, et al. Evaluating frontier models for dangerous capabilities. *arXiv preprint arXiv:2403.13793*, 2024.